



The National University of Computer and
Emerging Sciences

Introduction to Neural Networks

Machine Learning for Data Science

Dr. Akhtar Jamil

Department of Computer Science

Goals

- Review of Previous Lecture
- Today's Lecture
 - Neural Network and Backpropagation Desirable Features

Previous Lecture


Neural Network

- A popular method for the training of multilayer perceptron is the **back-propagation algorithm (Generalized Delta Rule)**
- *Forward phase:*
 - The input data is **propagated through the network**, layer by layer, until it reaches the output.
- *Backward phase:*
 - An **error is calculated** by comparing the output of the network with a desired response.
 - This **error is propagated** through the network in **backward direction**.

Neural Network

- **Hidden Neurons**
- The hidden neurons act as *feature detectors*
- The hidden neurons gradually “discover” the main **features** to understand the training data.
- They perform **nonlinear transformations** on the input
 - *Feature space.*
- Classes may be more **easily separated** from each other

Backpropagation algorithm



Backpropagation algorithm

- We can calculate the desired error for the output units as:

$$J(\mathbf{w}) \equiv 1/2 \sum_{k=1}^c (t_k - z_k)^2 = 1/2(\mathbf{t} - \mathbf{z})^2$$

- The **backpropagation learning rule** is based on **gradient descent**.
- The weights are initialized with random values, and are changed in a direction that will reduce the error:

$$\Delta \mathbf{w} = -\eta \frac{\partial J}{\partial \mathbf{w}}$$

- where η is the **learning rate**

Backpropagation algorithm

- The weight vector will be updated in iterative manner as:

$$w(m + 1) = w(m) + \Delta w(m)$$

$$w(m + 1) = w(m) - \eta \frac{\partial J}{\partial w}$$

- The weights throughout the network can be updated using.

$$\Delta w_{mn} = -\eta \frac{\partial J}{\partial w_{mn}}$$

- Now the challenge is to update the weights
 - Some are **not explicitly dependent on incoming weights**


Chain Rule

Give $y = g(u)$ and $u = h(x)$

Chain

Rule

$$\frac{dy_i}{dx_k} = \sum_{j=1}^J \frac{dy_i}{du_j} \frac{du_j}{dx_k}, \quad \forall i, k$$



Backpropagation algorithm

- Consider first the hidden-to-output weights, w_{jk} .
- we must use the chain rule for differentiation:

$$\frac{\partial J}{\partial w_{kj}} = \frac{\partial J}{\partial net_k} \frac{\partial net_k}{\partial w_{kj}} = \delta_k \frac{\partial net_k}{\partial w_{kj}}$$

- *Sensitivity of output unit* k is defined to be

$$\delta_k \equiv -\partial J / \partial net_k$$

Backpropagation algorithm

- Differentiate the **cost function with unit's net activation** to find how the overall error changes

$$J(\mathbf{w}) \equiv 1/2 \sum_{k=1}^c (t_k - z_k)^2$$

$$\delta_k \equiv -\partial J / \partial net_k = -\frac{\partial J}{\partial z_k} \frac{\partial z_k}{\partial net_k} = (t_k - z_k) f'(net_k)$$

- From the equation

$$\frac{\partial net_k}{\partial w_{kj}} = y_i \quad net_k = \sum_{j=1}^{n_H} y_j w_{kj} + w_{k0}$$


Backpropagation algorithm

- The learning rule for the **input-to-hidden** units can be calculated as:

$$\frac{\partial J}{\partial w_{ji}} = \frac{\partial J}{\partial y_j} \frac{\partial y_j}{\partial net_j} \frac{\partial net_j}{\partial w_{ji}}$$

Backpropagation algorithm

- The sensitivity at a hidden unit is proportional to the **weighted sum of the sensitivities at the output units**.
- The output unit sensitivities are thus **propagated “back” to the hidden units**.



Today's Lecture

Stochastic backpropagation

Algorithm 1 (Stochastic backpropagation)

```
1 begin initialize network topology (# hidden units),  $\mathbf{w}$ , criterion  $\theta, \eta, m \leftarrow 0$   
2   do  $m \leftarrow m + 1$   
3      $\mathbf{x}^m \leftarrow$  randomly chosen pattern  
4      $w_{ij} \leftarrow w_{ij} + \eta \delta_j x_i; \quad w_{jk} \leftarrow w_{jk} + \eta \delta_k y_j$   
5   until  $\nabla J(\mathbf{w}) < \theta$   
6 return  $\mathbf{w}$   
7 end
```

Batch backpropagation


Algorithm 2 (Batch backpropagation)

```
1 begin initialize network topology (# hidden units),  $\mathbf{w}$ , criterion  $\theta, \eta, r \leftarrow 0$ 
2   do  $r \leftarrow r + 1$  (increment epoch)
3      $m \leftarrow 0; \Delta w_{ij} \leftarrow 0; \Delta w_{jk} \leftarrow 0$ 
4     do  $m \leftarrow m + 1$ 
5        $\mathbf{x}^m \leftarrow$  select pattern
6        $\Delta w_{ij} \leftarrow \Delta w_{ij} + \eta \delta_j x_i; \Delta w_{jk} \leftarrow \Delta w_{jk} + \eta \delta_k y_j$ 
7     until  $m = n$ 
8      $w_{ij} \leftarrow w_{ij} + \Delta w_{ij}; w_{jk} \leftarrow w_{jk} + \Delta w_{jk}$ 
9   until  $\nabla J(\mathbf{w}) < \theta$ 
10 return  $\mathbf{w}$ 
11 end
```


Batch backpropagation

- *Stopping criterion:*
 - It is important to stop learning when **optimal values** of weight are obtained
- Two possibilities:
 - Define **Maximum number of epochs**
 - Reach a **minimum error threshold**
- Epochs
 - The **number of times the full training set** is fed to the model
 - One epoch means model has seen the full training set once.


Batch backpropagation




Error surfaces

- Generally, a **convex error surface** is desirable
- **Local minima**: if many local minima are present in the error surface, then it is **unlikely** that the network will find **the global minimum**.
- Another issue is **the presence of plateaus** — regions where the **error varies only slightly** as a function of weights.
 - **Training will be slow.**


Error surfaces



Error surfaces




Error surfaces



One Hot Encoding

- Machine learning models expect **numeric data for training**
 - e.g. preprocess text to some numeric representation
 - Image pixels are numeric, can extract more features
 - Sensor data can be represented in numeric form
- In case of classification, **the labels should be integers.**
- Instead of integers, **one-hot encoding** is more useful for training some neural networks

One Hot Encoding

Index	Animal		Index	Dog	Cat	Sheep	Lion	Horse
0	Dog	One-Hot code 	0	1	0	0	0	0
1	Cat		1	0	1	0	0	0
2	Sheep		2	0	0	1	0	0
3	Horse		3	0	0	0	0	1
4	Lion		4	0	0	0	1	0

Outputs as probabilities

- One approach is to choose the output neuron's nonlinearity to be **exponential rather than sigmoidal**:

$$z_k = \frac{e^{net_k}}{\sum_{m=1}^c e^{net_m}}$$

- Train the model with **one hot encoding**
- This is the **softmax method** — a smoothed or **softmax continuous version** of a *winner-take-all* nonlinearity.
- **Winner-take- all**, in which the maximum output is transformed to 1.0, and all others reduced to 0.0.


Activation Functions

- Backpropagation can work with **any activation function**
- The activation function should at least meet the criteria of
 - Differentiability, continuity, nonlinearity.
- The ***sigmoid function*** possesses the desirable properties of the activation function

Activation Functions


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU


$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Number of hidden units

- The **number of hidden neurons** (n_H) is crucial for extracting **powerful features**
 - The **complexity of the decision boundary**.
- For **linearly separable** data **few hidden units** are needed
- For **complex and nonlinear data**, **more hidden neurons** are needed.
- **No recommended number of hidden units**
- We should **not have more weights than the total number of training samples**

Number of hidden units



Initializing weights

- The weight values should be initialized to have fast and uniform learning.
 - All weights reach their optimum values at about the same time.
- If some weights converge significantly earlier than others (non-uniform learning) then the network may not perform well throughout the full range of inputs or for each class
- Example: when class C_i is learned well before C_j .
- A possible solution is to use data standardization

Initializing weights

- Setting weights to zero is **not a good idea**
- Choose weights randomly from a **single distribution to help insure uniform learning**.
- For standardized **input with d features**, we can randomly initialize the weight between **input – hidden** layers as:

$$-1/\sqrt{d} < w_{ji} < +1/\sqrt{d}$$


- Similarly for **hidden-output weights** can be randomly initialized as:

$$-1/\sqrt{n_H} < w_{kj} < +1/\sqrt{n_H}$$

Learning rates

- The **optimal learning rate** is the one which leads to the **minimum local error** in one
- Generally, learning step of $\eta = 0.1$ is good for starting
 - **Lower** it if the **cost function diverges**
 - **Raised** it the cost function is **too slow to converge**.
- How to identify slow convergence?
 - **Loss is gradually reducing** but take too many epochs

Learning rates




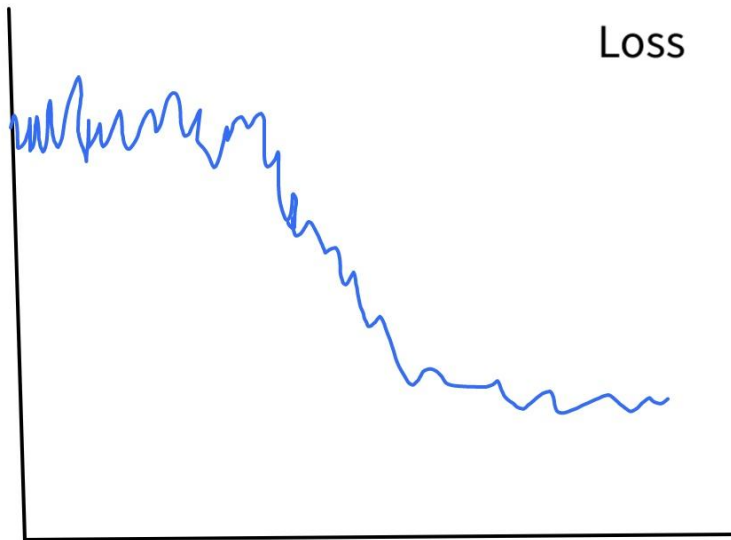
Momentum

- **Error** surfaces often have **plateaus regions**
 - The **slope** $dJ(\mathbf{w})/d\mathbf{w}$ is very small
 - Too many weights.
- **Momentum**- allows the network to **learn weights more quickly** when plateaus in the error surface exist.
- The **learning rule in backpropagation** can be updated as:

$$\mathbf{w}(m+1) = \underbrace{\mathbf{w}(m) + \Delta\mathbf{w}(m)}_{\text{gradient descent}} + \underbrace{\alpha\Delta\mathbf{w}(m-1)}_{\text{momentum}}$$

- α must be less than 1.0 for stability, typical value $\alpha = 0.9$.
- Affect is like **averaging the stochastic variations** in weight updates
- Can also help **overcome local minima**

Momentum



Exploding Gradients

- **Exploding gradients** are a problem where large error gradients accumulate
- Result in **very large weight updates** during training.
- Makes model **unstable and unable to learn** from your training data.
- Cause **poor predication** results

Weight decay

- Weight decay is a **regularization technique**
- Simplifying a network and avoiding overfitting is to impose a heuristic that the weights should be small.
- To prevent overfitting.
- To keep the weights small and avoid exploding gradient.
- This will help keep the weights as small as possible, preventing the weights to grow out of control, and thus avoid exploding gradient.

Weight decay

- Small weights favor models
- Popular due to its simplicity.
- After each weight update every weight is simply “decayed” or shrunk according to:

$$w^{\text{new}} = w^{\text{old}} (1 - \epsilon)$$

- where $0 < \epsilon < 1$

$$J_{ef} = J(\mathbf{w}) + \frac{2\epsilon}{\eta} \mathbf{w}^t \mathbf{w}$$

- Achieves a balance between error and overall weight.

Loss functions

- The squared error criterion is the most common training criterion
- It is simple to compute, non-negative,
 - Log loss
- another criterion function is based on the *Minkowski Minkowski error*:

$$J(\mathbf{w})_{ce} = \sum_{m=1}^n \sum_{k=1}^c t_{mk} \ln(t_{mk}/z_{mk})$$

References

- Chapter 4, Neural Networks and Learning Machines, Haykin
- Chapter 5, Pattern Recognition and Machine Learning, Bishop

Thank You 😊